

Machine Learning and Data Mining

Lecture 0.1: Introduction, Class Logistics, and Sources of Error



CS 434

1



Who is this guy?
Me in my natural habitat
What I study
Work toward the rolling and right, past the stairs. Walk into the living room and turn right. Stop by the end table.
Pick up the light green cup

Where I'm from
What I do for fun
His home is:

CS 434

2

CS434 – ML + DM

Today's Learning Objectives



Be able to answer:

- What is machine learning (ML)?
- When is it appropriate to use?
- How is this class structured?
- How will I be assessed / graded?
- How can I categorize different ML techniques?
- What is the standard ML formalism?
- Why must we make assumptions to learn?
- What are sources of error in ML?
- What is the idea behind k-nearest neighbors?

CS 434

3

What is Machine Learning?

"Field of study that gives computers the ability to learn without being explicitly programmed"
- Arthur Samuel

"Any change in a system that allows it to perform better the second time on repetition of the same task or on another task drawn from the same distribution"
- Herbert A. Simon

"Algorithms that automatically detect patterns in data and use the uncovered patterns to predict future data or other outcomes of interest"
- Kevin Murphy

"Algorithms that improve their performance (P) at some task (T) with experience (E)"
- Tom Mitchell



CS 434

4

What is Machine Learning? (Arthur Samuel's Answer)

Traditional Programming



Machine Learning



"Field of study that gives computers the ability to learn without being explicitly programmed"
- Arthur Samuel

CS 434

5

What is Machine Learning? (Tom Mitchell's Answer)


Machine Learning

Task T → Learning Algorithm ← Experience E (Data) → Performance P


"Algorithms that improve their performance (P) at some task (T) with experience (E)"
- Tom Mitchell

CS 434


6




7




8



9



10



11

- Where is Machine Learning used?
- Machine learning is already the preferred approach to:**
 - Speech recognition
 - Natural Language Processing
 - Computer Vision
 - Robotic Control
 - Recommender Systems
 - ...
 - This trend is growing**
 - Improved machine learning algorithms
 - Increased data capture and availability
 - Increased computing power
 - Increased demand for personalization to users and environments
- CS 434 © Stanford University [12]

When do we need learning?

CS 434

© Student Lens

13

When do we need learning?

1) Situations where humans can perform a task but can't describe how programmatically.

Image Classification

CS 434

© Student Lens

14

13

When do we need learning?

1) Situations where humans can perform a task but can't describe how programmatically.

Video Segmentation

CS 434

© Student Lens

15

When do we need learning?

1) Situations where humans can perform a task but can't describe how programmatically.

Speech Recognition

CS 434

© Student Lens

16

15

When do we need learning?

1) Situations where humans can perform a task but can't describe how programmatically.

This message seems dangerous
Similar messages were used to steal people's personal information. Avoid clicking links, downloading attachments, or replying with personal information.

[Learn more](#)

(This email originated from outside of UCalgary. Use caution with links and attachments.)

Dear Staff Employee,

Alberta Health Services (AHS) has introduced the Alert System to manage and minimize the risk of COVID-19. All staff are required to Logon COVID-19 Alert System to see the tasks and objectives.

To access, Click on: [COVID-19 Alert System](#).

Take care and keep safe.

Sincerely,
IT Helpdesk
©2020 All rights reserved

Dwayne Vincent
System Administrator
The Royal University
910 50 Street, Edmonton, Alberta
T6C 2G4
PH: 780.492.4000

SPAM

Spam Detection

CS 434

© Student Lens

17

When do we need learning?

2) Situations where the desired function is different for each individual user

Recommendation Systems

CS 434

© Student Lens

18

17

When do we need learning?

2) Situations where the desired function is different for each individual user

Ad Targeting

Ad Targeting

CS 434

19

When do we need learning?

2) Situations where the desired function is different for each individual user

Can I ask who is speaking so I can better understand your voice in the future?

Personalization

CS 434

20

When do we need learning?

3) Situations where human experts do not have sufficient knowledge (or time)

Drug / Material Discovery

Based on molecular structure, predict effect.

CS 434

21

When do we need learning?

4) Situations where the function of interest is changing rapidly.

Anticipating Demand

CS 434

22

When do we need learning?

4) Situations where the function of interest is changing rapidly.

Predicting Stock Markets

CS 434

23

Recent Public Successes of Machine Learning


AlphaGo
00:08:32
BBC NEWS

LEE SEDOL
00:00:27

nature
ALL SYSTEMS GO

CS 434

24



25



26



27



28




29



30

Recent Public Failures of Machine Learning



"we found empirical evidence for the existence of demographic differentials in the majority of the face recognition algorithms we studied"

- Patrick Grother, Lead author of NIST Study on Facial Recognition

Tested facial recognition products from >90 companies. For many, found false-identification rate **100x** higher for individuals of African or Asian descent than for Caucasians.

Many groups lobby for banning facial recognition technology in public places.

CS 434

31

31

Recent Public Failures of Machine Learning



CS 434

© Stefan Lee

© Stefan Lee

© Stefan Lee

32

Question Break!



Someone approaches you with a business proposal. They want you to build a machine learning system to predict whether a student will pass a class or not based on their transcripts. But due to FERPA, they have no data from previous classes and lack the budget to procure any data. Also, to protect their innovative idea, they won't let you ask for donated data publicly.

They are offering you 40% of the business equity - which they **insist** will be worth a fortune right after you finish. Obviously, no pay for now.

A Accept the job and start planning how big of a yacht you will buy with your share.

B Decline the job and avoid talking to this person again in the future.

CS 434

© Stefan Lee

33

33

CS434 – ML + DM

Today's Learning Objectives



Be able to answer:

- What is machine learning (ML)?
- When is it appropriate to use?
- How is this class structured?
- How will I be assessed / graded?
- How can I categorize different ML techniques?
- What is the standard ML formalism?
- Why must we make assumptions to learn?
- What are sources of error in ML?
- What is k-nearest neighbors?

© Stefan Lee

34

34

CS434 - Instructional Team

Instructor: Stefan Lee



Class Period: MW 4-5:50pm in COVL 216 (You're here. Well done. Kudos)

Office Hours: M 11:30-1:00pm on Zoom or KEC3013

Teaching Assistant Team (Office hours TBD)



Eric
Slyman



Xiangxi
Shi



Jing
Wang



Ian
Tassin

CS 434

© Stefan Lee

35

35

CS434 - What do you need to do for points?

Grading / Assessment:

- **Quiz Participation** (10%) - Regular conceptual quizzes / surveys. Graded on completion -- where completion means making a reasonable effort at responding to the questions, judged at the TA's discretion. Always due on Sundays.
- **Assignments** (50%) - Five assignments over the term each due on Canvas. First one out today and due next Wednesday.
- **Exams** (40%) - One midterm and a final - 20% each. In-class. Open book / note.

© Stefan Lee

36

36

CS434 - Don't be late or cheat.

Late Work Policy:

Late assignments will be accepted within 48 hours of the deadline with a penalty. You will receive 95% of your regular grade for late submission submitted within 24 hours of the deadline. Late submission submitted between 24 and 48 hours past the deadline will receive 90% of your regular grade. No submission will be accepted after 48 hours. No late quizzes or exams will be accepted.

Collaboration Policy:

Students are expected to work independently. High-level discussion of topics is allowed, but any low-level details, sharing of code, or collaborative programming is prohibited. Likewise, the use of code from online resources is prohibited. You do not want to get caught doing this.

Only Exception: Work however closely as you want for homework 0.

CS 434

© Student Lee

37

37

CS434 - Rough Schedule

Week 0	Introduction and Course Logistics	HW0 Out.
Week 1	K-Nearest Neighbor and MLE/MAP	HW0 Due. HW1 Out.
Week 2	MAP (cont.) and Linear Regression	HW1 Due. HW2 Out.
Week 3	Polynomial Regression and Logistic Regression	
Week 4	Perceptron and SVM	
Week 5	SVM (cont.) and Kernel Methods	HW2 Due.
Week 6	Midterm (Review + Exam)	Midterm Exam. HW3 Out.
Week 7	Multiclass Classifiers and Neural Networks	
Week 8	Decision Trees and Ensemble Methods	HW3 Due. HW4 Out.
Week 9	Ensemble Methods (cont.) and Clustering	
Week 10	Dimensionality Reduction and Final Review	HW4 Due.
Finals Week		Final Exam.
Forever after	Bask in the glory of your newfound machine learning knowledge	

CS 434

Schedule subject to change

© Student Lee

38

38

CS434 - Background Material

Am I well prepared for this course?

Maybe? Are you comfortable with most of these (or willing to learn):

- Manipulating probabilities and distributions [Probability]
- Taking derivatives and reasoning about integrals [Calculus]
- Performing common matrix operations [Linear Algebra]
- Standard computer science concepts like...
 - Programming (skeleton code will be in Python)
 - Computational complexity (big O notation)
- Handling abstract mathematical concepts [Computer Science]
- Proving things by manipulating equations [Discrete Math]

Homework 0 is out today. Not worth many points and graded on effort but should give you a sense for if you'll need to review some things for this class.

CS 434

© Student Lee

39

39

CS434 - Reaching for Success

How to do well in this class?

Come to class on time. Pay attention. Take notes. Ask questions during lecture or afterwards via Discord or in office hours. Rewatch recordings for parts you missed.

Start your homework early. You will always have a homework to do but we don't plan for it to actually take the full two weeks you have to do it. That said, don't wait till the last day – you will very probably not do well on it.

If you don't know background material but are still keen to learn ML, ask for help and study hard.

Just like saying "Eat healthy, sleep eight hours, avoid stress, and exercise regularly", I know this is obvious advice. Following it is the hard part.

CS 434

© Student Lee

40

40

CS434 - Machine Learning and Data Mining

What this class is:

An introduction to the fundamentals of machine learning. You will do math to derive and understand learning algorithms. You will implement learning algorithms from scratch and apply them to provided datasets.

What this class isn't:

A tutorial on how to do machine learning in framework X, Y, or Z. Also, this is not a deep learning class.


CS 434

© Student Lee

41

41

CS434 - Day-to-Day



42

CS434 - Day-to-Day

110 Minutes

50 Minutes
1-5 min admin

50 Minutes
5-10 min break

CS 434

43

Question Break!

If homework zero seems hard to me,
Professor Stefan is telling me that I should:

A Drop the class and give up on my dreams of understanding machine learning 😞

B View the recorded review sessions in areas you find challenging and ask questions in Discord

C Consider whether the extra work of refreshing or learning these skills fits into my life

D Invest all my money in Dogecoin How could you think this is what I was saying? So wrong. Much not listening. (This is not financial advice).

CS 434

44

CS434 - Machine Learning and Data Mining

It is okay not to know something

Imposter Syndrome

Reality

If you don't recognize a symbol or notation I use, let me know. I'm happy to explain it and then you will - easy fix.

CS 434

45

CS434 – ML + DM

Today's Learning Objectives

Be able to answer:

- What is machine learning (ML)?
- When is it appropriate to use?
- How is this class structured?
- How will I be assessed / graded?
- How can I categorize different ML techniques?
- What is the standard ML formalism?
- Why must we make assumptions to learn?
- What are sources of error in ML?
- What is the idea behind k-nearest neighbors?

CS 434

46

Machine Learning In a Nutshell

- **Thousands of machine learning algorithms**
 - Hundreds new every year
- **Decades of ML research oversimplified:**
 - All of Machine Learning:
 - Learn a mapping from input to output $g: X \rightarrow Y$
 - X: emails, Y: {spam, notspam}

CS 434

47

Machine Learning In a Nutshell

Input: x	(images, text, emails...)
Output: y	(spam or non-spam...)
(Unknown) Target Function $f: X \rightarrow Y$	(the "true" mapping / reality)
Data $(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)$	(our observations of the world)
Model / Hypothesis Class $\mathcal{H} = \{g: X \rightarrow Y\}$	(the space of possible models)

CS 434



48

Machine Learning Topics

Supervised Learning

Learn to predict output from input given example (x, y) pairs

Regression: Predicting a (or multiple) continuous values as output





CS 434

Machine Learning Topics

Supervised Learning

Learn to predict output from input given example (x, y) pairs

Classification: Predicting a (or multiple) discrete variable as output



CS 434

50


49

Machine Learning Topics

Unsupervised Learning

Only given a set of input instances (x)

Clustering: Discover self-similar groups within the data



Example: Find groups of houses with similar properties


CS 434

Machine Learning Topics

Unsupervised Learning

Only given a set of input instances (x)

Density Estimation: Estimate the underlying distribution generating our data



Example: Estimate how likely is a given house with these properties

51


52

Machine Learning Topics

Unsupervised Learning

Only given a set of input instances (x)

Dimensionality Reduction: Represent high-dim data as low-dim data



Example: Finding a 2D subspace in a 3D feature space


CS 434

53

Machine Learning Topics

Reinforcement Learning

Multiple interactions with an environment yield a reward signal. No optimal outputs are provided. Goal: learn a policy that optimizes rewards



Barely going to touch this one.

CS 434

54

No Assumptions → No Learning

It's clear we are interested in mappings from input to output. Let's get concrete.

 Willow (legally distinct from Zillow)

Your records contain d attributes for each house...

- Lets assume binary: HasPool? >300sqft?, IsHaunted?...
- Implies $x \in \{0,1\}^d$

... and whether it sold within a month.

- Implies $y \in \{0,1\}$

How many possible mappings from x to y are there?

CS 434 © Stanford University [55]

55

No Assumptions → No Learning

How many possible mappings $g: \{0,1\}^d \rightarrow \{0,1\}$ from x to y are there?

What is a mapping? For each unique x , provide an output y .

x	y	x	y	x	y
0 0 ... 0	0	0 0 ... 0	1	0 0 ... 0	1
0 0 ... 1	1	0 0 ... 1	1	0 0 ... 1	0
.
.
1 1 ... 1	0	1 1 ... 1	0	1 1 ... 1	0

...

Making no assumptions about the function -- 2^{2^d} possible!

CS 434 © Stanford University [56]

56

No Assumptions → No Learning

Will data save us?

Each unique data point fills in one of these rows for us. Say we have n .

x	y
0 0 ... 0	1
0 0 ... 1	
.	
.	
1 1 ... 1	0

Still 2^{2^d-n} possible mappings. Far, far too many.

We need to make some sort of assumption about the relationship between input and output.

CS 434 © Stanford University [57]

57

Question Break!

You want to predict shoe size from someone's height. You are given a dataset of height/shoe-size pairs.

This is an example of a:

A supervised classification problem B unsupervised clustering problem

C supervised regression problem D unsupervised density estimation problem

CS 434 © Stanford University [58]

58

Question Break!

You want to estimate how likely it is for someone to have a certain shoe size / height combo. You are given a dataset of height/shoe-size pairs.

This is an example of a:

A supervised classification problem B unsupervised clustering problem

C supervised regression problem D unsupervised density estimation problem

CS 434 © Stanford University [59]

59

Question Break!

You want to group people by their shoe-size and height. You are given a dataset of height/shoe-size pairs.


This is an example of a:

A supervised classification problem B unsupervised clustering problem


C supervised regression problem D unsupervised density estimation problem

CS 434 © Stanford University [60]


60




61




62




63




64



65



66



67

Question Break!

Green points are our observed data, which fit do you think is best?

H₁: A straight line fit to the data.
H₂: A parabolic fit to the data.
H₈: An 8-degree polynomial fit to the data.


A Left One (First-degree polynomial fit)
B Middle One (2nd-degree polynomial fit)
C Right One (8th-degree polynomial fit)
D I wasn't paying attention and the quiet woke me up. What are we talking about again?

CS 434 © Duda et al.


68




69




70




71




72




73




74




75




76




77




13




79




80




81



82



83



84

Question Break!

Which source(s) of error tend(s) towards zero with infinite data?

A Optimization Error **B Estimation Error**
C Modelling Error **D Bayes Error**

CS 434

85


Question Break!

Which source(s) of error tend(s) towards zero as computation goes to infinity?

A Optimization Error **B Estimation Error**
C Modelling Error **D Bayes Error**

CS 434

86



87

What is Bayes Error?

Irreducible error inherit in the function being approximated - nothing we can fix.

Bayes Error

CS 434

88

What is Bayes Error?

Bayes Error
Irreducible error inherit in the function being approximated - nothing we can fix.

Bayes Optimal Classifier:

Suppose we know the true distribution $P(Y|X)$ and for each x we encounter we predict:

$$\hat{y} = \operatorname{argmax}_y P(Y=y|X=x)$$

How often would we be wrong?

- It depends on how random $P(y|x)$ is. Consider a binary problem $y \in \{0,1\}$
 - If $P(y=0|X=x) = 0.6 \rightarrow$ Wrong 40% of the time for x .
- Could we do better? Nope. Proof can be found here: https://en.wikipedia.org/wiki/Bayes_classifier

CS 434

89

CS434 – ML + DM


Today's Learning Objectives

Be able to answer:


- What is machine learning (ML)?
- When is it appropriate to use?
- How is this class structured?
- How will I be assessed / graded?
- How can I categorize different ML techniques?
- What is the standard ML formalism?
- Why must we make assumptions to learn?
- What are sources of error in ML?
- What is the idea behind k-nearest neighbors?

CS 434


90




91




92



93




94




95




96



97




98



99




100



101



102



CS 434 Adapted from Dhruv Batra

© Sudar Lee

Hays and Efros, SIGGRAPH 2007

103



CS 434 Adapted from Dhruv Batra

© Sudar Lee

Hays and Efros, SIGGRAPH 2007

104



CS 434 Adapted from Dhruv Batra

© Sudar Lee

Hays and Efros, SIGGRAPH 2007

105



CS 434

© Sudar Lee

106